

# Unit 1

## Simple Linear Regression

Prof. Phil Schniter



THE OHIO STATE UNIVERSITY

ECE 4300: Introduction to Machine Learning, Sp20

# Learning objectives

- Understand how to load datasets in Python using pandas
- Visualize data using a scatter plot
- Understand sample mean, variance, and covariance
- Describe a linear model for data
  - Identify the target variable and predictor
- Compute the least-squares fit using the regression formula
- Compute the  $R^2$  measure-of-fit
- Visually assess goodness-of-fit and identify causes of poor fit

# Outline

- Motivating Example: Predicting Automobile MPG
- Linear Model
- Least-Squares Fit: The Problem
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit

## Example: Predicting automobile mpg

- We will consider the task of predicting the fuel efficiency (in mpg) of a car from features like # cylinders, horsepower, weight, etc.
- Demo in Jupyter notebook: `demo01_auto_mpg.ipynb`
  - Learn from demo, and test your understanding in the lab
- Uses data from UCI library: <https://archive.ics.uci.edu/ml/>
- Data is loaded using Python's `pandas` library
  - Pandas routines have many options!
  - Learn from examples; Google is your friend!
  - Pandas `read_csv` command creates a `dataframe` object with 3 main components:
    - `df.values`: numerical values in `Numpy` format
    - `df.columns`: column labels
    - `df.index`: row labels

# Result of pandas' read\_csv

```
import pandas as pd
import numpy as np
```

```
In [3]: names = ['mpg', 'cylinders', 'displacement', 'horsepower',
                 'weight', 'acceleration', 'model year', 'origin', 'car name']
```


```
In [4]: df = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/'+
                        'auto-mpg/auto-mpg.data',
                        header=None, delim_whitespace=True, names=names, na_values='?')
df.head(6)
```

```
Out[4]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18.0	8	307.0	130.0	3504.0	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693.0	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150.0	3436.0	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150.0	3433.0	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140.0	3449.0	10.5	70	1	ford torino
5	15.0	8	429.0	198.0	4341.0	10.0	70	1	ford galaxie 500

# Visualizing the data

- When possible, visualize the data before working with it.
- Python has MATLAB-like plotting features in the `matplotlib` module.



## Visualizing the Data

We load the `matplotlib` module to plot the data in MATLAB.

```
In [15]: import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

First, let's extract some data columns and create arrays:




```
In [16]: xstr = 'displacement'
x = np.array(df[xstr])
y = np.array(df['mpg'])
```

Now we can create a scatter plot:

```
In [17]: plt.plot(x,y,'o')
plt.xlabel(xstr)
plt.ylabel('mpg')
plt.grid(True)
```

# Postulating a model

- What relationships do you see?
- Is there a mathematical model relating the variables?
- How well can you predict mpg from these variables?




# Outline

- Motivating Example: Predicting Automobile MPG
- **Linear Model**
- Least-Squares Fit: The Problem
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit



# Describing and visualizing data

- $y$ : the variable we are trying to predict
  - Many names: **target**, **regressand**, **response variable**, **dependent variable**, **label**
- $x$ : the variable we are using for prediction
  - Many names: **feature**, **predictor**, **regressor**, **attribute**, **independent variable**
  - For now we consider a single variable, but in general there can be many
- Data: the set of pairs  $\{(x_i, y_i)\}_{i=1}^n$ 
  - Each pair is called a **sample**
  - We will use  $n$  for the number of samples
- A **scatter plot** is used to visualize the data pairs




# Single-variable linear model

- Assume a linear relationship:

$$y \approx \beta_0 + \beta_1 x$$

- $\beta_1$ : slope
- $\beta_0$ : intercept
- $\beta = [\beta_0, \beta_1]^T$  are the **parameters** of the model. Also called **coefficients** or **weights**
- Why this model?
  - easy to interpret & analyze
  - simple to optimize parameters (as we will see)
- When is this a good model?



The **regression line** is shown in red

# Outline

- Motivating Example: Predicting Automobile MPG
- Linear Model
- **Least-Squares Fit: The Problem**
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit

# Linear model: The residual

- Note:  $x$  does not *exactly* predict  $y$ :


$$y \approx \beta_0 + \beta_1 x$$

- Let's model this behavior explicitly using a **residual**,  $\epsilon$ :

$$y = \beta_0 + \beta_1 x + \epsilon$$

- For the  $i$ th sample, we then have:

- **predicted value**:  $\hat{y}_i = \beta_0 + \beta_1 x_i$
- **residual**:  $\epsilon_i = y_i - \hat{y}_i$



The **residual**  $\epsilon_i$  is the vertical deviation from  $y_i$  to the regression line

# Least-squares fit

- How do we choose the model parameters  $\beta = [\beta_0, \beta_1]^T$ ?
- Minimize the **residual sum of squares** (RSS):

$$\text{RSS}(\beta_0, \beta_1) \triangleq \sum_{i=1}^n \epsilon_i^2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

[https://en.wikipedia.org/wiki/Residual\\_sum\\_of\\_squares](https://en.wikipedia.org/wiki/Residual_sum_of_squares)

- Also called the **sum of squared errors** (SSE) and **sum of square residuals** (SSR)
- Note that  $\epsilon_i$  and  $\hat{y}_i$  are implicitly functions of  $[\beta_0, \beta_1]$
- The value of  $\beta$  that minimizes RSS is the **least-squares fit**.
  - Geometrically: minimizes sum of squared distances to the regression line.

# The optimization approach: A general ML recipe

## General ML problem

- Assume a **model** with some **parameters**
- Get data
- Choose a **loss function**
- Find parameters that **minimize** loss

## Simple Linear Regression

→ Linear model:  $\hat{y} = \beta_0 + \beta_1 x$

→ Data:  $\{(x_i, y_i)\}_{i=1}^n$

→  $\text{RSS}(\beta_0, \beta_1) \triangleq \sum_{i=1}^n (y_i - \hat{y}_i)^2$

→ Find  $(\beta_0, \beta_1)$  that minimizes  $\text{RSS}(\beta_0, \beta_1)$

# Outline

- Motivating Example: Predicting Automobile MPG
- Linear Model
- Least-Squares Fit: The Problem
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit

# Sample mean, variance, & standard deviation

## ■ Sample mean:

$$\bar{x} \triangleq \frac{1}{n} \sum_{i=1}^n x_i, \quad \bar{y} \triangleq \frac{1}{n} \sum_{i=1}^n y_i$$


## ■ Sample variance:

$$s_x^2 \triangleq \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2, \quad s_y^2 \triangleq \frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2$$

- Note: some authors use  $\frac{1}{n-1}$  to give an unbiased estimate. More on this later.

## ■ Sample standard deviation (SD): $s_x, s_y$

- Simply the square-root of the sample variance




<https://en.wikipedia.org/wiki/Variance>



# Visualizing sample mean & SD on scatter plot

- Sample means  $\bar{x}$  and  $\bar{y}$ :
  - The center of mass in each axis
- Standard deviations  $s_x$  and  $s_y$ :
  - The “spread” in each axis about the mean
  - If the data was Gaussian distributed...
    - 68% of points < 1 SD from mean
    - 95% of points < 2 SDs from mean
    - 99.7% of points < 3 SDs from mean
- What are your estimates of  $\bar{x}$ ,  $\bar{y}$ ,  $s_x$ ,  $s_y$  from the above scatter plot?




# Computing the sample mean & SD in Python

- We can exactly compute  $\bar{x}, \bar{y}, s_x, s_y$  using the **Numpy** package in **Python**

```
In [27]: xm = np.mean(x)
          ym = np.mean(y)
          sxx = np.mean((x-xm)**2)
          syy = np.mean((y-ym)**2)
```

```
xm      = 104.47,      ym= 23.45
sqrt(sxx)= 38.44,    sqrt(syy)= 7.80
```



# Sample covariance

- The **sample covariance** is


$$s_{xy} \triangleq \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$$

which indicates how “related”  $\{x_i\}$  and  $\{y_i\}$  are.

- The value  $s_{xy}$  is easier to interpret after normalization. This is done via the **sample (Pearson) correlation coefficient**:

$$\rho_{xy} \triangleq \frac{s_{xy}}{s_x s_y} \in [-1, 1]$$

- The property  $\rho_{xy} \in [-1, 1]$  is a consequence of the Cauchy-Schwarz inequality.
- Example scatterplots for datasets with various  $\rho_{xy}$ :



[https://en.wikipedia.org/wiki/Pearson\\_correlation\\_coefficient](https://en.wikipedia.org/wiki/Pearson_correlation_coefficient)

# Alternative expressions for sample variance & covariance

- Recall that the sample variance was defined as  $s_x^2 \triangleq \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$
- A very useful alternative formula can be found by expanding the square:

$$s_x^2 = \frac{1}{n} \sum_{i=1}^n x_i^2 - 2\bar{x} \frac{1}{n} \sum_{i=1}^n x_i + \bar{x}^2 = \frac{1}{n} \sum_{i=1}^n x_i^2 - \bar{x}^2 \quad \Rightarrow \quad \frac{1}{n} \sum_{i=1}^n x_i^2 = s_x^2 + \bar{x}^2$$

- Similarly, for the sample covariance we had  $s_{xy} \triangleq \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$
- and a useful alternative expression is

$$s_{xy} = \frac{1}{n} \sum_{i=1}^n x_i y_i - \bar{x} \bar{y} \quad \Rightarrow \quad \frac{1}{n} \sum_{i=1}^n x_i y_i = s_{xy} + \bar{x} \bar{y}$$

# Notation

- We will use the following notation in this class (we'll try to be consistent)
- Note: some books/authors use different notations

Statistic	Notation	Formula	Python
sample mean	$\bar{x}$	$\frac{1}{n} \sum_{i=1}^n x_i$	xm
sample variance	$s_x^2 = s_{xx}$	$\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$	sxx
sample standard deviation	$s_x = \sqrt{s_{xx}}$	$\sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2}$	sx
sample covariance	$s_{xy}$	$\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$	sxy
sample correlation coefficient	$\rho_{xy}$	$\frac{s_{xy}}{s_x s_y}$	rhoxy

# Outline

- Motivating Example: Predicting Automobile MPG
- Linear Model
- Least-Squares Fit: The Problem
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit

# Minimizing RSS



- To minimize  $\text{RSS}(\beta_0, \beta_1)$ , we find the  $\beta_0$  and  $\beta_1$  that zero the partial derivatives

$$\frac{\partial \text{RSS}(\beta_0, \beta_1)}{\partial \beta_0} = 0, \quad \frac{\partial \text{RSS}(\beta_0, \beta_1)}{\partial \beta_1} = 0$$

Because RSS is quadratic with a positive Hessian, the zero-gradient point is guaranteed to occur at a minimum (more on this later)

- After some manipulations (see next page), we obtain the optimal values

$$\beta_1 = \frac{s_{xy}}{s_{xx}} = \frac{\rho_{xy}s_y}{s_x}, \quad \beta_0 = \bar{y} - \beta_1\bar{x}$$



# Minimizing RSS: Derivation

The minimum RSS is achieved by values of  $(\beta_0, \beta_1)$  that zero the gradient, i.e.,

$$0 = \frac{\partial \text{RSS}(\beta_0, \beta_1)}{\partial \beta_0} = \frac{\partial}{\partial \beta_0} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2 = -2 \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i) \quad (1)$$

$$0 = \frac{\partial \text{RSS}(\beta_0, \beta_1)}{\partial \beta_1} = \frac{\partial}{\partial \beta_1} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2 = -2 \sum_{i=1}^n x_i (y_i - \beta_0 - \beta_1 x_i) \quad (2)$$

Starting with (1), we can multiply both sides by  $-\frac{1}{2n}$  to give

$$0 = \frac{1}{n} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i) = \underbrace{\frac{1}{n} \sum_{i=1}^n y_i}_{\bar{y}} - \beta_0 - \beta_1 \underbrace{\frac{1}{n} \sum_{i=1}^n x_i}_{\bar{x}} \Leftrightarrow \boxed{\beta_0 = \bar{y} - \beta_1 \bar{x}} \quad (3)$$

Doing the same with (2) gives

$$0 = \frac{1}{n} \sum_{i=1}^n x_i (y_i - \beta_0 - \beta_1 x_i) = \frac{1}{n} \sum_{i=1}^n x_i y_i - \bar{x} \beta_0 - \beta_1 \frac{1}{n} \sum_{i=1}^n x_i^2 \quad (4)$$



# Minimizing RSS: Derivation (continued)

Plugging in (3) into (4) gives

$$0 = \underbrace{\frac{1}{n} \sum_{i=1}^n x_i y_i - \bar{x}\bar{y}}_{s_{xy}} - \beta_1 \underbrace{\left( \frac{1}{n} \sum_{i=1}^n x_i^2 - \bar{x}^2 \right)}_{s_{xx}} \Leftrightarrow \boxed{\beta_1 = \frac{s_{xy}}{s_{xx}}}$$

To find the minimum RSS, we plug the optimal value of  $\beta_0$  into the RSS definition to get


$$\begin{aligned} \text{RSS}(\beta_0, \beta_1) &= \sum_i (y_i - \beta_0 - \beta_1 x_i)^2 = \sum_i \left( (y_i - \bar{y}) - \beta_1 (x_i - \bar{x}) \right)^2 \\ &= \underbrace{\sum_i (y_i - \bar{y})^2}_{ns_{yy}} - 2\beta_1 \underbrace{\sum_i (y_i - \bar{y})(x_i - \bar{x})}_{ns_{xy}} + \beta_1^2 \underbrace{\sum_i (x_i - \bar{x})^2}_{ns_{xx}} \end{aligned} \quad (5)$$

and then we plug the optimal value of  $\beta_1$  into (5) to get

$$\begin{aligned} \text{RSS}(\beta_0, \beta_1) &= n \left( s_{yy} - 2 \frac{s_{xy}^2}{s_{xx}} + \frac{s_{xy}^2}{s_{xx}} \right) = n \left( s_{yy} - \frac{s_{xy}^2}{s_{xx}} \right) \\ &= n \left( 1 - \frac{s_{xy}^2}{s_{xx} s_{yy}} \right) s_{yy} = n(1 - \rho_{xy}^2) s_{yy} \end{aligned}$$

# Automobile demo

- Applied to our Python demo...



```

xm = np.mean(x)
ym = np.mean(y)
sxx = np.mean((x-xm)**2)
syy = np.mean((y-ym)**2)
syx = np.mean((y-ym)*(x-xm))
beta1 = syx/sxx
beta0 = ym - beta1*xm

```

beta0=39.94, beta1=-0.16

$$\widehat{\text{mpg}} = \beta_0 + \beta_1 \times \text{horsepower}$$

- Another good simple-linear-regression demo can be found at <https://stattrek.com/regression/regression-example.aspx?Tutorial=AP>

# Outline

- Motivating Example: Predicting Automobile MPG
- Linear Model
- Least-Squares Fit: The Problem
- Sample Mean, Variance, Standard Deviation, Covariance
- Least-Squares Fit: The Solution
- Assessing Goodness-of-Fit

# $R^2$ Goodness-of-fit

- Key question: How good is this linear prediction?
- Let's split the variance-of- $y$  into two parts:

$$s_y^2 = \underbrace{\left[ s_y^2 - \frac{\text{RSS}}{n} \right]}_{\text{explained by } x} + \underbrace{\left[ \frac{\text{RSS}}{n} \right]}_{\text{unexplained by } x} = \left( \underbrace{\left[ 1 - \frac{\text{RSS} / n}{s_y^2} \right]}_{\text{explained by } x} + \underbrace{\left[ \frac{\text{RSS} / n}{s_y^2} \right]}_{\text{unexplained by } x} \right) s_y^2$$

- The *fraction* of the variance-of- $y$  explained by  $x$  is

$$1 - \frac{\text{RSS} / n}{s_y^2} \triangleq R^2,$$

known as the “coefficient of determination”


[https://en.wikipedia.org/wiki/Coefficient\\_of\\_determination](https://en.wikipedia.org/wiki/Coefficient_of_determination)

- Note that  $R^2 \in [0, 1]$
- Ex:  $R^2 = 0.48$  means that “48% of  $s_y^2$  is explained by  $x$ .”
- With LS coefficients  $(\beta_0, \beta_1)$ , we have

$$\text{RSS}(\beta_0, \beta_1) = n(1 - \rho_{xy}^2) s_y^2 \quad \Rightarrow \quad R^2 = \rho_{xy}^2$$

# Visualizing the correlation coefficient

- $R^2 = \rho_{xy}^2 \approx 1$ : Linear model is a very good fit
- $R^2 = \rho_{xy}^2 \approx 0$ : Linear model is a very poor fit
- $\beta_1 = \frac{\rho_{xy}s_y}{s_x} \Rightarrow \text{sgn}(\beta_1) = \text{sgn}(\rho_{xy})$



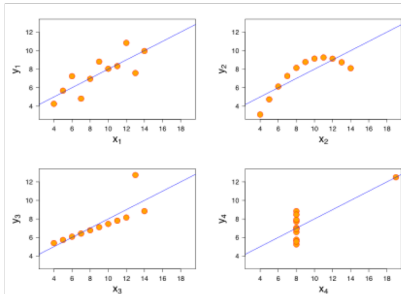
← linear models with  
varying  $\rho_{xy}$  (and  
 $s_x = s_y$ )

← linear models with  
 $\rho_{xy} \in \{-1, 0, 1\}$   
(and varying  $s_y$ )

← no linear  
relationship

# If the prediction error is large. . .



- There are many sources of error in linear models
- Example on right:
  - All 4 datasets have same regression line
  - Their prediction errors are large for different reasons
- A lot can be learned from visually inspecting the scatter plot!



# A better model for the automobile example?

- What if we predicted the inverse:  

$$\frac{1}{\text{mpg}} \approx \beta_0 + \beta_1 \times \text{horsepower}$$
- This uses a nonlinear data transformation followed by linear regression
- Will explore this idea later in the course



# Learning objectives

- Understand how to load datasets in Python using pandas
- Visualize data using a scatter plot
- Understand sample mean, variance, and covariance
- Describe a linear model for data
  - Identify the target variable and predictor
- Compute the least-squares fit using the regression formula
- Compute the  $R^2$  measure-of-fit
- Visually assess goodness-of-fit and identify causes of poor fit