

# Unit 4

## Feature Selection and LASSO

Prof. Phil Schniter



THE OHIO STATE UNIVERSITY

ECE 4300: Introduction to Machine Learning, Sp20

# Learning objectives


- Understand motivation and idea behind **feature selection**
- Understand feature selection methods based on:
  - **exhaustive search**
  - **stepwise selection**
  - **target cross-correlation**
  - **regularization**
- Understand **ridge regression** and **LASSO**:
  - interpret their **coefficient paths**
  - implement LASSO using **sklearn**
  - know how to select the **regularization strength** using cross-validation
- Understand connections to **ML estimation** and **MAP estimation**

# Outline

- Motivating Example: Predicting Prostate Cancer
- Feature Selection
- Ridge Regression and LASSO
- Probabilistic Interpretations of Regularized Regression
- Extension to Vector-Valued Targets

# Prostate-specific antigen (PSA) testing

- High PSA is linked to prostate cancer
  - PSA levels are easily monitored
  - A common tool for screening
- Classic 1989 study by Thomas et al:
  - Measured PSA level of 102 men prior to prostate removal
  - Measured various biometrics
  - Biometrics include cancer volume, prostate weight, age, etc.
- Machine-learning problem:
  - Can we predict PSA from these biometrics?



Stamey, et al., "Prostate specific antigen in the diagnosis and treatment of adenocarcinoma of the prostate. II. Radical prostatectomy treated patients," *The Journal of Urology*, 141.5 (1989): 1076-1083.

# PSA Dataset

- Prostate dataset widely used in ML classes
- Can be downloaded from many websites
- Data samples from 97 patients
- 8 features, shown on right
- Target variable = *lpsa* (log PSA)

```
# Get data
url = 'https://web.stanford.edu/~hastie/ElemStatLearn/datasets/prostate.data'
df = pd.read_csv(url, sep='\t', header=0)
df = df.drop('Unnamed: 0', axis=1) # skip the column of indices
```

The data frame has the following components:

```
lcavol      log(cancer volume)
lweight     log(prostate weight)
age         age
lbph        log(benign prostatic hyperplasia amount)
svi         seminal vesicle invasion
lcp         log(capsular penetration)
gleason     Gleason score
pgg45       percentage Gleason scores 4 or 5
lpsa        log(prostate specific antigen)
```

# First attempt: Linear Regression

- Let's try first with multiple linear regression:

$$y \approx \hat{y} = \beta_0 + \beta_1 x_1 + \cdots + \beta_d x_d$$

- $y = \text{lpsa}$  (target PSA level)
  - $\{x_j\}_{j=1}^d$  are biometric features with  $d = 8$
- Why linear regression?

- Coefficients are easy to fit (via LS)
  - Coefficients are easy to interpret
    - larger  $|\beta_j|$  means  $x_j$  has larger effect on PSA

```
import sklearn.model_selection

# construct leave-one-out-cross-val object
loocv = sklearn.model_selection.KFold(n_splits=nsamp)

# construct linear regression model
linreg = linear_model.LinearRegression()
```

```
from sklearn.model_selection import cross_val_score

scores = cross_val_score(linreg, X, y, cv=loocv, scoring='neg_mean_squared_error')
print(-np.mean(scores))

0.4104899877500149
```

- Can we do better?

# Outline

- Motivating Example: Predicting Prostate Cancer
- Feature Selection
- Ridge Regression and LASSO
- Probabilistic Interpretations of Regularized Regression
- Extension to Vector-Valued Targets

# Feature Selection

- From last lecture:
  - Too many features  $\Rightarrow$  large error variance
  - This motivates using fewer features. But which ones?
  - **Feature selection**: use only the best *subset* of  $d$  total features
- Feature selection via **exhaustive search**:
  - Main idea: use K-fold CV to test *every possible* subset
  - This is the **optimal** approach to feature selection
  - But, with large  $d$ , testing  $2^d$  subsets may be computationally impractical!
- Suboptimal feature selection methods:
  - **Stepwise selection**
  - **Correlation-based** methods
  - **Regularization-based** methods, e.g., LASSO



# Stepwise selection (or stepwise regression)

## ■ Forward selection

- First use CV to find the **single** feature yielding the lowest RSS
- Then, **add one** of the remaining features so that the pair provides the lowest RSS
- Repeat until RSS starts to increase, or maximum allowed # features is reached

## ■ Backwards elimination

- First use **all**  $d$  features and compute the RSS using CV
- Then, **remove one** of the features so that the remaining features give lowest RSS
- Repeat until RSS starts to increase, or only one feature remains

## ■ These methods are suboptimal

- They look only one step ahead (i.e., “greedy”)
- A better approach is LASSO, discussed later

# Feature selection via cross-correlation

Here is another heuristic strategy. Again, assume  $d$  total features.

- **Maximize cross-correlation with target:**

- For each  $p \in \{1, \dots, d\}$ , find the  $p$  features that are most correlated with the target (i.e., compute  $r_{yx_j} = \frac{1}{n} \sum_{i=1}^n y_i x_{ij}$  for each  $j$ , and choose the  $p$  values of  $j$  giving largest  $|r_{yx_j}|$ )
- Use cross-validation to optimize the model-order  $p$

- Is this a good idea?

- Not necessarily
- Two features might both be highly correlated with the target, but provide redundant information, in which case only one of them should be used
- There exist more sophisticated versions that penalize correlations among features  
[https://en.wikipedia.org/wiki/Feature\\_selection#Correlation\\_feature\\_selection](https://en.wikipedia.org/wiki/Feature_selection#Correlation_feature_selection)

- A better approach is LASSO, discussed next

# Outline

- Motivating Example: Predicting Prostate Cancer
- Feature Selection
- Ridge Regression and LASSO
- Probabilistic Interpretations of Regularized Regression
- Extension to Vector-Valued Targets

# Regularization

- Previously, we optimized the linear regression coefficients  $\beta$  via LS:

$$\hat{\beta} = \arg \min_{\beta} \text{RSS}(\beta) \quad \text{for} \quad \text{RSS}(\beta) = \sum_{i=1}^n (y_i - \hat{y}_i(\beta))^2$$

- Can we modify this to perform **feature selection**?
- Idea: **Penalize** the use of each feature (i.e., penalize  $\beta_j \neq 0$ )
  - If the penalty results in  $\hat{\beta}_j = 0$ , then the  $j$ th feature is not used
  - In particular, add a “**regularization term**”  $\phi(\beta)$  to the optimization objective:

$$\hat{\beta} = \arg \min_{\beta} J(\beta) \quad \text{for} \quad J(\beta) = \text{RSS}(\beta) + \phi(\beta)$$

for some  $\phi(\cdot)$  that encourages  $\hat{\beta}_j = 0$  for non-informative features  $j$

- How should we choose  $\phi(\cdot)$ ?

# L1 and L2 regularization


- L2 regularization:

$$\phi(\boldsymbol{\beta}) = \alpha \sum_{j=1}^d |\beta_j|^2$$

- L1 regularization:

$$\phi(\boldsymbol{\beta}) = \alpha \sum_{j=1}^d |\beta_j|$$

- Both penalize  $\beta_j \neq 0$ , but in different ways
- The overall strength of regularization is controlled by  $\alpha \geq 0$
- Note: regularization does *not* involve the intercept term  $\beta_0$ , since we do *not* want to penalize it!



# Data standardization

## ■ Motivation:

- The L1 and L2 regularizers penalize all coefficients  $\beta_j$  *uniformly*
- But if some  $x_j$  are much bigger than others, then some  $\beta_j$  may be much bigger than others, in which case  $\beta_j$  should not be treated uniformly
- Can avoid this issue by **normalizing** the sizes (i.e., sample variances) of  $x_j$
- If we also **remove the mean** of  $y$  and each  $x_j$ , then conveniently we get  $\beta_0 = 0$

## ■ Procedure:

- “**Standardize**” target and features to have **sample mean 0** and **sample variance 1**:

$$x_{ij} \leftarrow (x_{ij} - \bar{x}_j) / s_{x_j} \quad \text{and} \quad y_i \leftarrow (y_i - \bar{y}) / s_y$$

- Design a predictor  $\beta$  using the standardized data
- When test data  $x$  arrives, either standardize  $x$  using the above quantities, or un-standardize the predictor  $\beta$  (see homework)

# Ridge regression and LASSO

Assuming  $(\mathbf{y}, \mathbf{X})$  has been standardized, and setting  $\boldsymbol{\beta} = [\beta_1, \dots, \beta_d]^\top$ ,

- **Ridge regression** cost function:

$$J_{\text{ridge}}(\boldsymbol{\beta}) = \underbrace{\sum_{i=1}^n (y_i - \hat{y}(\boldsymbol{\beta}))^2}_{=\text{RSS}(\boldsymbol{\beta})} + \underbrace{\alpha \sum_{j=1}^d |\beta_j|^2}_{\text{L2 regularization}} = \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|^2 + \alpha \|\boldsymbol{\beta}\|^2$$

- **LASSO** cost function:

$$J_{\text{lasso}}(\boldsymbol{\beta}) = \underbrace{\sum_{i=1}^n (y_i - \hat{y}(\boldsymbol{\beta}))^2}_{=\text{RSS}(\boldsymbol{\beta})} + \underbrace{\alpha \sum_{j=1}^d |\beta_j|}_{\text{L1 regularization}} = \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|^2 + \alpha \|\boldsymbol{\beta}\|_1$$

Note:

- $\|\boldsymbol{\beta}\| = \|\boldsymbol{\beta}\|_2$  is known as the “**L2 norm**” or Euclidean norm
- $\|\boldsymbol{\beta}\|_1$  is known as the “**L1 norm**” or “taxi-cab” norm or “Manhattan” norm

# Ridge regression

- Recall the ridge cost:

$$J_{\text{ridge}}(\beta) = \text{RSS}(\beta) + \alpha \sum_{j=1}^d |\beta_j|^2 = \|\mathbf{y} - \mathbf{X}\beta\|^2 + \alpha \|\beta\|^2$$

- Similar to the derivation of  $\beta_{\text{ls}}$ , can show that

$$\beta_{\text{ridge}} \triangleq \arg \min_{\beta} J_{\text{ridge}}(\beta) = (\mathbf{X}^T \mathbf{X} + \alpha \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

- Why use the penalty  $\|\beta\|^2$ ?

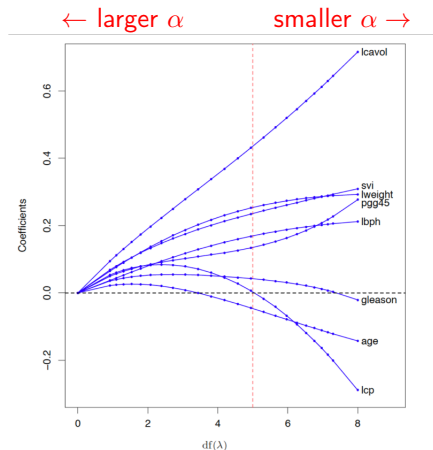
- When **columns of  $\mathbf{X}$  are correlated**, the unregularized LS solution  $\beta_{\text{ls}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$  can have very large values due to the inverse
- Equivalently, feature correlation can make the cost surface  $\|\mathbf{y} - \mathbf{X}\beta\|^2$  very stretched and its minimum,  $\beta_{\text{ls}}$ , very far from the origin
- Problem: large  $\beta_{\text{ls}}$  implies  $\hat{y} = \beta_{\text{ls}}^T x$  is sensitive to test data  $x$  (i.e., overfitting)
- By penalizing  $\|\beta\|^2$ , we discourage large  $\beta$  and thus help to **reduce overfitting**



# Coefficient path of ridge regression

- The “**coefficient path**” is the plot of all  $\beta_j$  versus the regularization strength  $\alpha$
- With ridge regression, larger  $\alpha$  leads to smaller  $\beta_j$  but *not fewer* non-zero  $\beta_j$
- Choose  $\alpha$  via cross-validation

Figure from book: Hastie, Tibshirani, Friedman,  
*The Elements of Statistical Learning*



**FIGURE 3.8.** Profiles of ridge coefficients for the prostate cancer example, as the tuning parameter  $\lambda$  is varied. Coefficients are plotted versus  $df(\lambda)$ , the effective degrees of freedom. A vertical line is drawn at  $df = 5.0$ , the value chosen by cross-validation.

# LASSO


- Recall LASSO cost:

$$J_{\text{lasso}}(\beta) = \text{RSS}(\beta) + \alpha \sum_{j=1}^d |\beta_j|_1 = \|\mathbf{y} - \mathbf{X}\beta\|^2 + \alpha \|\beta\|_1$$

- No closed-form expression for  $\beta_{\text{lasso}} \triangleq \arg \min_{\beta} J_{\text{lasso}}(\beta)$ 
  - but convex optimization problem  $\Rightarrow$  tractable solution
  - many fast numerical solvers: FISTA, ADMM, glmnet
  - implemented in `sklearn` via the `Lasso` method
- Why use penalty  $\|\beta\|_1$ ?
  - Leads to **exactly zero**  $\beta_j$ , and thus **feature selection**!
    - $\alpha$  controls # of nonzero  $\beta_j$
  - Careful: the non-zero  $\beta_j$  are biased towards 0
    - discard them and keep only the *indices* of informative features,  $\{j : \beta_j \neq 0\}$
    - then do standard linear regression (i.e., LS) with the informative  $x_j$


# Coefficient path of LASSO

- The “**coefficient path**” is the plot of all  $\beta_j$  versus the regularization strength  $\alpha$
- With LASSO, larger  $\alpha$  leads to fewer non-zero  $\beta_j$
- The LASSO path suggests which features  $\{x_j\}$  are most informative
  - In the PSA demo, LASSO suggests `lcavol` is most informative
- Choose  $\alpha$  via cross-validation



# Summary of ridge regression and LASSO

- LASSO (L1 penalty)
  - Tends to produce many exactly-zero  $\beta_j$
  - Great for feature selection!
  - But no closed-form solution: solve numerically
- Ridge regression (L2 penalty)
  - Can write solution in closed-form
  - Tends to produce many small  $\beta_j$
  - Not useful for feature selection
  - But helps with correlated features



# Implementing LASSO with sklearn

- **sklearn** has a **Lasso** method
- On right, we choose the regularization weight  $\alpha$  using K-fold cross-validation:
  - Outer loop over folds  $k$
  - Inner loop over a **grid of  $\alpha$**
  - First compute  $RSS_{k,\alpha}$  for all  $k$  and  $\alpha$
  - Then compute  $\overline{RSS}_\alpha$  and  $SE_\alpha$  for each  $\alpha$

```
# Create a k-fold cross validation object
nfold = 10
kf = sklearn.model_selection.KFold(n_splits=nfold,shuffle=True)

# Create the LASSO model. We use the `warm start` parameter so
# This speeds up the fitting.
model = linear_model.Lasso(warm_start=True)

# Regularization values to test
nalpha = 100
alphas = np.logspace(-3,1,nalpha)

# MSE for each alpha and fold value
mse = np.zeros((nalpha,nfold))
for ifold, ind in enumerate(kf.split(X)):

    # Get the training data in the split
    Itr,Its = ind
    X_tr = X[Itr,:]
    y_tr = y[Itr]
    X_ts = X[Its,:]
    y_ts = y[Its]

    # Compute the lasso path for the split
    for ia, a in enumerate(alphas):

        # Fit the model on the training data
        model.alpha = a
        model.fit(X_tr,y_tr)

        # Compute the prediction error on the test data
        y_ts_pred = model.predict(X_ts)
        mse[ia,ifold] = np.mean((y_ts_pred-y_ts)**2)

# Compute the mean and standard deviation over the different fo
mse_mean = np.mean(mse,axis=1)
mse_std = np.std(mse,axis=1) / np.sqrt(nfold-1)
```


# Applying the one-standard-error rule

Select  $\alpha$  via one-standard-error rule:

- Find  $\alpha_{\min} = \arg \min_{\alpha} \overline{RSS}_{\alpha}$
- Set  $\overline{RSS}_{\text{tgt}} = \overline{RSS}_{\alpha_{\min}} + SE_{\alpha_{\min}}$
- Find simplest model (largest  $\alpha$ ) such that  $\overline{RSS}_{\alpha} < \overline{RSS}_{\text{tgt}}$

```
# Find the minimum MSE and MSE target
imin = np.argmin(mse_mean)
alpha_min = alphas[imin]
mse_min = mse_mean[imin]
mse_tgt = mse_min + mse_std[imin]

# Find the least complex model with mse_mean < mse_tgt
I = np.where(mse_mean < mse_tgt)[0]
iopt = I[-1]
alpha_opt = alphas[iopt]
print("Optimal alpha = %f" % alpha_opt)
```



# Using LASSO for feature selection

- Having found  $\alpha$ , we next compute the LASSO coefficients  $\beta_{\text{lasso}}$  on the full training data
  - LASSO selects only 3 features for prediction!
- We then **isolate these features** and use them for LS-based linear regression
  - Note  $\beta_{\text{lasso}}$  is used only for feature selection
  - Why? Non-zeros in  $\beta_{\text{lasso}}$  are biased

```
model.alpha = alpha_opt
model.fit(X,y)

# Print the coefficients
for i, c in enumerate(model.coef_):
    print("%8s %f" % (names_x[i], c))
```

```
lcavol 0.457465
lweight 0.103180
age 0.000000
lbph 0.000000
svi 0.120609
lcp 0.000000
gleason 0.000000
```

```
X1=np.zeros((nsamp, 3))
X1[:,0]=X[:,0]
X1[:,1]=X[:,1]
X1[:,2]=X[:,4]

scores = cross_val_score(linreg, X1, y, cv=loo cv, scoring='neg_mean_squared_error')
print(-np.mean(scores))
```

0.3985076811922328 ← the RSS improved!


# Outline

- Motivating Example: Predicting Prostate Cancer
- Feature Selection
- Ridge Regression and LASSO
- Probabilistic Interpretations of Regularized Regression
- Extension to Vector-Valued Targets



# Estimation of statistical parameters

- At its core, machine *learning* is essentially about **estimating statistical parameters** in a probabilistic model of the data
- We now describe the **maximum-likelihood (ML)** and **maximum a posteriori (MAP)** approaches to this important problem
- Say  $y$  is a random variable that depends on some statistical parameters  $\theta$ .
  - Then  $y$  is fully described by its **probability density function (pdf)**  $p(y|\theta)$ .
  - To visualize the pdf, we usually plot  $p(y|\theta)$  versus  $y$  for some hypothesized  $\theta$
- Example:  $y \sim \mathcal{N}(\mu, \sigma^2)$ , i.e.,  $y$  is **Gaussian** with mean  $\mu$  and variance  $\sigma^2$ 
  - Then  $p(y|\theta) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp(-\frac{1}{2\sigma^2}(y - \mu)^2)$  with parameters  $\theta = [\mu, \sigma^2]^\top$
  - This pdf is a “bell curve” centered at  $\mu$  with width  $\sigma$ :



# Maximum-likelihood (ML) estimation

- Now say that we observe independent samples  $\mathbf{y} = [y_1, \dots, y_n]^T$  of a random variable  $y$  with pdf  $p(y|\boldsymbol{\theta})$ .
- For this fixed  $\mathbf{y}$ , the function  $p(\mathbf{y}|\boldsymbol{\theta})$  versus  $\boldsymbol{\theta}$  is called the “**likelihood function**”
- Since we’ve assumed  $\{y_i\}_{i=1}^n$  are **independent & identically distributed (i.i.d.)**,

$$p(\mathbf{y}|\boldsymbol{\theta}) = \prod_{i=1}^n p(y_i|\boldsymbol{\theta})$$

- A common way of estimating  $\boldsymbol{\theta}$  from observed  $\mathbf{y}$  is to **maximize the likelihood**:

$$\boldsymbol{\theta}_{\text{ml}} \triangleq \arg \max_{\boldsymbol{\theta}} p(\mathbf{y}|\boldsymbol{\theta}) = \arg \max_{\boldsymbol{\theta}} \{\ln p(\mathbf{y}|\boldsymbol{\theta})\} = \arg \min_{\boldsymbol{\theta}} \{-\ln p(\mathbf{y}|\boldsymbol{\theta})\}$$

where  $\ln(\cdot)$  and negation are often used to simplify the expression

- ML estimation uses no prior belief about  $\boldsymbol{\theta}$ ; it only fits the observed data

# ML estimation for linear regression

- Can we use ML estimation to **fit the parameters  $\beta$  of linear regression?** Yes!
- Suppose that our data obeys the linear-Gaussian model

$$\mathbf{y} = \mathbf{X}\beta + \epsilon \text{ with } \{\epsilon_i\} \sim \text{i.i.d. } \mathcal{N}(0, \sigma^2)$$

assuming standardized data, and thus  $\beta = [\beta_1, \dots, \beta_d]^\top$  (i.e., no intercept  $\beta_0$ )

- For this model, can show that the likelihood fxn becomes

$$p(\mathbf{y}|\mathbf{X}, \beta) = \frac{\exp(-\frac{1}{2\sigma^2} \|\mathbf{y} - \mathbf{X}\beta\|^2)}{(2\pi\sigma^2)^{n/2}}$$

- The **ML estimate** is then

$$\beta_{\text{ml}} = \arg \min_{\beta} \{ -\ln p(\mathbf{y}|\mathbf{X}, \beta) \} = \arg \min_{\beta} \|\mathbf{y} - \mathbf{X}\beta\|^2 = \beta_{\text{ls}}$$

- So, under this linear-Gaussian model, **the ML estimate is the least-squares fit!**

# Maximum a posteriori (MAP) estimation

- To incorporate a **prior belief** on  $\beta$ , we can use **MAP estimation**
- The MAP estimate of  $\beta$  from  $\mathbf{y}$  is

$$\beta_{\text{map}} \triangleq \arg \max_{\beta} p(\beta | \mathbf{X}, \mathbf{y})$$

i.e., the **most probable**  $\beta$  given the data  $(\mathbf{X}, \mathbf{y})$

- **Bayes rule** says

$$p(\beta | \mathbf{X}, \mathbf{y}) = \frac{p(\mathbf{y} | \mathbf{X}, \beta) p(\beta)}{p(\mathbf{y})} \text{ for "prior" pdf } p(\beta)$$

which implies that

$$\beta_{\text{map}} = \arg \max_{\beta} \frac{p(\mathbf{y} | \mathbf{X}, \beta) p(\beta)}{p(\mathbf{y})} = \arg \max_{\beta} \{p(\mathbf{y} | \mathbf{X}, \beta) p(\beta)\}$$

- Interpretations:

**likelihood**  $p(\mathbf{y} | \mathbf{X}, \beta)$ : how well  $\beta$  agrees with data  $(\mathbf{X}, \mathbf{y})$

**prior**  $p(\beta)$ : how well  $\beta$  agrees with prior belief

- Key point: MAP estimation uses both likelihood *and* prior

# Regularized linear regression is MAP estimation

- It's often simpler to formulate
 
$$\begin{aligned}\beta_{\text{map}} &= \arg \max_{\beta} \{p(\mathbf{y}|\mathbf{X}, \beta)p(\beta)\} \\ &= \arg \min_{\beta} \{ -\ln [p(\mathbf{y}|\mathbf{X}, \beta)p(\beta)] \} \\ &= \arg \min_{\beta} \{ -\ln p(\mathbf{y}|\mathbf{X}, \beta) - \ln p(\beta) \}\end{aligned}$$
- The linear-Gaussian model:  $\mathbf{y} = \mathbf{X}\beta + \epsilon$  with i.i.d.  $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$ , gives

$$-\ln p(\mathbf{y}|\mathbf{X}, \beta) = \frac{1}{2\sigma^2} \|\mathbf{y} - \mathbf{X}\beta\|^2 + \text{const}$$

and so

$$\begin{aligned}\beta_{\text{map}} &= \arg \min_{\beta} \left\{ \frac{1}{2\sigma^2} \|\mathbf{y} - \mathbf{X}\beta\|^2 - \ln p(\beta) \right\} \\ &= \arg \min_{\beta} \left\{ \underbrace{\|\mathbf{y} - \mathbf{X}\beta\|^2}_{\text{RSS}(\beta)} - \underbrace{2\sigma^2 \ln p(\beta)}_{\phi(\beta)} \right\}\end{aligned}$$

- Thus MAP estimation under the linear-Gaussian model is equivalent to **regularized linear regression**!

# MAP interpretations of Ridge Regression and LASSO

- Recall that

$$\beta_{\text{map}} = \arg \min_{\beta} \{ \|\mathbf{y} - \mathbf{X}\beta\|^2 - 2\sigma^2 \ln p(\beta) \}$$

- If our prior belief is that  $\beta_j$  is i.i.d.  $\mathcal{N}(0, v)$ , then

$$p(\beta) = \prod_{j=1}^d \frac{\exp(-\frac{1}{2v}(\beta_j)^2)}{\sqrt{2\pi v}} \Rightarrow \ln p(\beta) = -\frac{1}{2v}\|\beta\|^2 + \text{const}$$

$$\Rightarrow \beta_{\text{map}} = \arg \min_{\beta} \{ \|\mathbf{y} - \mathbf{X}\beta\|^2 + \frac{\sigma^2}{v}\|\beta\|^2 \} \Leftrightarrow \text{Ridge Regression, } \alpha = \frac{\sigma^2}{v}$$

- If our prior belief is that  $\beta_j$  is i.i.d. Laplacian(0,  $\lambda$ ), then

$$p(\beta) = \prod_{j=1}^d \frac{\exp(-\frac{1}{\lambda}|\beta_j|)}{2\lambda} \Rightarrow \ln p(\beta) = -\frac{1}{\lambda}\|\beta\|_1 + \text{const}$$

$$\Rightarrow \beta_{\text{map}} = \arg \min_{\beta} \{ \|\mathbf{y} - \mathbf{X}\beta\|^2 + \frac{2\sigma^2}{\lambda}\|\beta\|_1 \} \Leftrightarrow \text{LASSO, } \alpha = \frac{2\sigma^2}{\lambda}$$

# Outline

- Motivating Example: Predicting Prostate Cancer
- Feature Selection
- Ridge Regression and LASSO
- Probabilistic Interpretations of Regularized Regression
- Extension to Vector-Valued Targets

# Linear regression with vector-valued targets

- Until now we've focused on linearly predicting a **scalar-valued target**  $y_i$  from features  $\mathbf{x}_i \in \mathbb{R}^d$  using a **coefficient vector**  $\boldsymbol{\beta} \in \mathbb{R}^{d+1}$  with intercept term  $\beta_0$ :

$$y_i \approx [1 \quad \mathbf{x}_i^\top] \boldsymbol{\beta}$$

- We can extend this approach to **vector-valued targets**  $\mathbf{y}_i^\top \in \mathbb{R}^K$  as follows

$$\mathbf{y}_i^\top \triangleq [y_{i1} \quad \cdots \quad y_{iK}] \approx [1 \quad \mathbf{x}_i^\top] \underbrace{[\beta_1 \quad \cdots \quad \beta_K]}$$

where now we use a **coefficient matrix**  $\mathbf{B} \in \mathbb{R}^{(d+1) \times K}$ .  $\triangleq \mathbf{B}$

- Incorporating all of the training samples  $i = 1, \dots, n$ , we get the model

$$\underbrace{\begin{bmatrix} \mathbf{y}_1^\top \\ \vdots \\ \mathbf{y}_n^\top \end{bmatrix}}_{\triangleq \mathbf{Y}} = \begin{bmatrix} y_{11} & \cdots & y_{1K} \\ \vdots & & \vdots \\ y_{n1} & \cdots & y_{nK} \end{bmatrix} \approx \underbrace{\begin{bmatrix} 1 & \mathbf{x}_1^\top \\ \vdots & \vdots \\ 1 & \mathbf{x}_n^\top \end{bmatrix}}_{\mathbf{A}} \underbrace{[\beta_1 \quad \cdots \quad \beta_K]}_{\mathbf{B}}$$



# Regularized regression with vector-valued targets

- With a linear-Gaussian model and standardized data (no intercepts), we get

$$\mathbf{Y} = \mathbf{X}\mathbf{B} + \mathbf{E} \quad \text{with i.i.d. } \epsilon_{ik} \sim \mathcal{N}(0, \sigma^2)$$

- For **Ridge Regression**, we would solve for

$$\mathbf{B}_{\text{ridge}} \triangleq \arg \min_{\mathbf{B}} \{ \|\mathbf{Y} - \mathbf{X}\mathbf{B}\|_F^2 + \alpha \|\mathbf{B}\|_F^2 \}$$

where  $\|\mathbf{B}\|_F^2 \triangleq \sum_{j,k} b_{jk}^2$  is the (squared) **matrix Frobenius norm**

- We don't use the matrix norm  $\|\mathbf{B}\| = \|\mathbf{B}\|_2$ , which has a different meaning!
- Meanwhile, for **LASSO**, we would solve for

$$\mathbf{B}_{\text{lasso}} \triangleq \arg \min_{\mathbf{B}} \{ \|\mathbf{Y} - \mathbf{X}\mathbf{B}\|_F^2 + \alpha \|\mathbf{B}\|_1 \}$$

where  $\|\mathbf{B}\|_1 \triangleq \sum_{j,k} |b_{jk}|$  is the **matrix L1 norm**

- **sklearn**'s linear regression methods have no trouble with these extensions

# Learning objectives

- Understand motivation for and concept of **feature selection**
- Understand feature selection methods based on:
  - **exhaustive search**
  - **stepwise selection**
  - **target cross-correlation**
  - **regularization**
- Understand **ridge regression** and **LASSO**:
  - interpret their **coefficient paths**
  - implement LASSO using **sklearn**
  - know how to select the **regularization strength** using cross-validation
- Understand connections to **ML estimation** and **MAP estimation**
- Understand how to handle **vector-valued targets**