

# Unit 10

## Random Forests and other Ensemble Methods

Prof. Phil Schniter



THE OHIO STATE UNIVERSITY

ECE 4300: Introduction to Machine Learning, Sp20

# Learning objectives

- Understand intuition behind ensemble methods: “the wisdom of the crowd”
- Understand parallel ensemble methods
  - bagging, pasting
  - random feature selection
- Understand decision trees
  - feature thresholding and decision regions
  - training via top-down tree induction
  - homogeneity metrics: variance reduction, gini impurity
  - ensemble extension: random forests
- Understand boosting, or sequentially trained ensemble methods
  - Adaboost
  - gradient boosting
  - XGBoost

# Outline

- Parallel Ensemble Methods: Bagging and Pasting
- Decision Trees and Random Forests
- Boosting: Sequentially Trained Ensemble Methods

# The wisdom of the crowd

- Suppose that you want to determine the heavy side of a biased coin
  - Although you don't know it, suppose "heads" is 51% and "tails" is 49%
- To predict the heavy side, you flip it  $N$  times and choose the **majority vote**
  - You'll be correct with probability
 

$N$	1	11	101	1001	10001
prob	51%	53%	58%	74%	97.7%
  - These values follow from the cdf of the binomial distribution
  - Interpretation: As you use more independent trials, the variance decreases
- Implications for classification:
  - Say you have 10001 binary classifiers, each correct only 51% of the time
    - Alone, each does barely better than random guessing ("**weak learner**")
  - A majority vote will classify with 97.7% accuracy! ("**strong learner**")
  - Caveat: The classifiers must generate **independent** outcomes

# Voting classifiers

- Given an **ensemble** of base classifiers ...
  - the **hard voting classifier** takes the majority vote of hard decisions
  - the **soft voting classifier** averages their soft outputs (i.e., pmfs) and then chooses the maximizing class ... and usually works better
  - Both are implemented in `sklearn.VotingClassifier`
- The voting classifier will be better than the base classifiers if the base classifiers are sufficiently **diverse** (i.e., sufficiently independent)
- There are two ways to generate a diverse ensemble of classifiers:
  - 1 Train the classifiers with **different data**
  - 2 Make the classifiers **structurally different**:
    - Use different classification methods (e.g., LR, SVC, NN)
    - Use same method but add randomness (e.g., choose random features)
- Similar ideas apply to an ensemble of base regressors

# Bagging and pasting

- Two main ways to train an ensemble using data diversity:
  - 1 **pasting**: draw training samples  $i$  **without replacement**
    - each sample is used exactly once by one predictor
  - 2 **bagging** (or **bootstrap aggregating**): draw training samples  $i$  **with replacement**
    - each sample may be used several times, or never
- bagging usually works better than pasting
- Bagging:
  - Usually draw  $n$  samples ( $n$  is total # training samples)
  - For large  $n$ , this implies  $e^{-1} \approx 37\%$  samples go unused (per predictor)
  - Can use the “**out of bag**” samples for cross-validation!
- Implemented in `sklearn.BaggingClassifier` & `sklearn.BaggingRegressor`:
  - Choose `bootstrap=True` for bagging (or `=False` for pasting)
  - Choose `oob_score=True` to report out-of-bag cross-validation

# Random patches and random subspaces

- Using **random subsets of features** is another way to add diversity
  - Useful when features are high-dimensional and redundant
  - Can sample features with replacement (i.e., “bootstrap”) or without
- Terminology:
  - Using both random data and features is known as “**random patches**”
  - Using random features but full data is known as “**random subspaces**”
- Implemented in `sklearn.BaggingClassifier` & `sklearn.BaggingRegressor`:
  - Set `max_features < d` for feature randomization
  - Set `bootstrap_features=True` for bootstrapping



# Outline

- Parallel Ensemble Methods: Bagging and Pasting
- Decision Trees and Random Forests
- Boosting: Sequentially Trained Ensemble Methods





# Decision trees

- Consider a supervised learning task (e.g., regression or classification)
  - Suppose data is  $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$  with  $\mathbf{x}_i \in \mathbb{R}^d$  and  $y_i \in \mathbb{R}$
- Approach:
  - 1 Partition domain  $\mathbb{R}^d$  into  $L$  regions  $\{R_\ell\}_{\ell=1}^L$
  - 2 Output  $z_\ell \in \mathbb{R}$  whenever  $\mathbf{x} \in R_\ell$ 
    - For regression, could set  $z_\ell$  at sample mean of  $\{y_i\}_{i \in S_\ell}$ , where  $S_\ell = \{i : \mathbf{x}_i \in R_\ell\}$
    - For classification, could set  $z_\ell$  at sample mode, i.e., most common value in  $\{y_i\}_{i \in S_\ell}$
- Regression example with dimension  $d = 2$ :



# Why are they “trees”?

- Suppose the regions are constructed by thresholding one feature at a time
  - The decision boundaries are always parallel to the coordinate axes
- Then can view prediction as a **decision tree**:
  - The domain of  $\mathbf{x}$  (i.e.,  $\mathbb{R}^d$ ) forms the **root** of the tree
  - The decision regions  $\{R_\ell\}$  are the  $L$  **leaves** of the tree
- Example with dimension  $d = 2$  and  $L = 5$  leaves:



# Top-down induction of decision trees

Decision trees are trained in a **top-down** manner:

- Start with entire data set:  $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$  and define set  $S \triangleq \{1, \dots, n\}$
- First split  $S$  into subsets  $(S_1, S_2)$ , so that labels  $\{y_i\}$  are most **homogeneous** (i.e., similar) within each subset
  - Splitting is performed by thresholding some feature  $j$ :
    - Ordinal case:  $S_1 = \{i : x_{ij} \leq t\}$  and  $S_2 = \{i : x_{ij} > t\}$
    - Categorical case:  $S_1 = \{i : x_{ij} \in \{A, C\}\}$  and  $S_2 = \{i : x_{ij} \in \{B, D\}\}$
  - To find “best” split, must search jointly over feature  $j$  and threshold  $t$
  - Many ways to measure homogeneity (will discuss later)
- Then split each subset  $S_1$  and  $S_2$  further, using the same procedure
- Repeat until ...
  - labels are perfectly homogeneous within a subset (e.g., all  $y_i$  are same), or
  - stopping condition: e.g., subsets have min # samples, tree has max depth, etc


# Decision tree example

Example: Risk prediction with  $n = 6$  samples,  $d = 2$  features,  $L = 3$  leaves.  
 Feature  $j = 1$  (age) is ordinal and feature  $j = 2$  (car type) is categorical:

Tid	Age	Car Type	Class
0	23	Family	High
1	17	Sports	High
2	43	Sports	High
3	68	Family	Low
4	32	Truck	Low
5	20	Family	High

Numeric

Categorical



Age=40, CarType=Family  $\Rightarrow$  Class=Low

<https://web.fhnw.ch/personenseiten/taoufik.nouri/Data%20Mining/Course/Course3/DM-Part%203.htm>

# Homogeneity metrics for regression

- Recall that the goal is to *maximize homogeneity* within  $S_1$  and  $S_2$ 
  - Equivalently, we want to *minimize inhomogeneity* within  $S_1$  and  $S_2$
- Let us first consider **regression**, where labels  $y_i \in \mathbb{R}$
- Inhomogeneity could be measured by the **variance** after splitting  $S \rightarrow (S_1, S_2)$ :
  - The mean in subset  $S_\ell$  is  $\mu_\ell \triangleq \frac{1}{|S_\ell|} \sum_{i \in S_\ell} y_i$ , for  $\ell \in \{1, 2\}$
  - The variance in subset  $S_\ell$  is  $v_\ell \triangleq \frac{1}{|S_\ell|} \sum_{i \in S_\ell} (y_i - \mu_\ell)^2$ , for  $\ell \in \{1, 2\}$
  - Thus the (average) variance after splitting is  $v \triangleq \frac{|S_1|}{|S|} v_1 + \frac{|S_2|}{|S|} v_2$
- Another option is to minimize the **absolute error**  $v_\ell \triangleq \frac{1}{|S_\ell|} \sum_{i \in S_\ell} |y_i - \mu_\ell|$
- We'll see a more sophisticated metric when discussing XGBoost on page 31

# Homogeneity metrics for classification

- Now consider  $K$ -ary **classification**, where  $y_i \in \{a_1, \dots, a_K\} \forall i$ 
  - For subset  $S_\ell$ , the **empirical label pmf** is  $\{p_{\ell k}\}_{k=1}^K$ , where  $p_{\ell k} = \frac{\sum_{i \in S_\ell} \mathbb{1}_{y_i = a_k}}{|S_\ell|}$
- Again, want to *maximize homogeneity*, or minimize inhomogeneity, in  $S_1$  &  $S_2$
- One popular measure of *inhomogeneity* is **Gini impurity**:

$$I_G = \frac{|S_1|}{|S|} I_{G,1} + \frac{|S_2|}{|S|} I_{G,2}, \quad \text{where } I_{G,\ell} \triangleq 1 - \sum_{k=1}^K p_{\ell k}^2$$

- Note that  $I_{G,\ell} \in [0, 1 - \frac{1}{K}]$
- $I_{G,\ell} = 0$  means perfectly pure/homogeneous, i.e., constant  $y_i$  for  $i \in S_\ell$
- Other criteria include **entropy**, **misclassification error**, **chi-square**
  - See **Hastie, Tibshirani, Friedman, *The Elements of Statistical Learning*, 2nd Ed., 2009**

# Advantages and disadvantages of decision trees

## Advantages:

- Very **general**
  - Very few assumptions made on data; can work with any dataset
  - Don't need to standardize, but should “balance” the dataset
- Very **interpretable**: “white box model”
  - Easy to understand how a tree arrives at its prediction
- Prediction is **fast**:  $O(\log_2 n)$  decisions to process a test sample  $x$


## Disadvantages:

- Training can be **expensive** if all features  $j$  are considered at every split
  - Solution: Restrict to a few randomly chosen features  $j$  at each split
- Prone to **overfitting**
  - Highly dependent on training data: changing one sample can change entire tree!
  - Can regularize by enforcing min # samples per subset, max depth, etc

~> There is a good overview of decision trees in the [sklearn documentation](#)!

# Random forests

- Decision trees tend to **overfit**: low bias & high variance
- Idea: Use a **random forest**: A random collection of trees generated by ...
  - using a random training subset to construct each tree (i.e., bagging or pasting)
  - using a random subset of features  $j$  per *split*
- As with other ensemble methods, results are averaged to make a prediction
- Implemented in `sklearn.RandomForestClassifier` & `Regressor`



from [machine-learning-algorithms-ensemble-methods-bagging-boosting-and-random-forests](#)




# Outline

- Parallel Ensemble Methods: Bagging and Pasting
- Decision Trees and Random Forests
- Boosting: Sequentially Trained Ensemble Methods

# Boosting

- Previously we described ensemble methods that train **in parallel**
- Now we discuss ensemble methods that train **in series**
  - This is called “**boosting**”
- Some of the most popular boosting methods are:
  - adaptive boosting, or **Adaboost** (1996)
  - **gradient boosting** (2001)
  - extreme gradient boosting or **XGBoost** (2016)



# Adaboost

- **Adaboost** trains an ensemble of predictors  $\{f_m(\cdot)\}_{m=1}^M$  sequentially as follows:

For  $m = 1 \dots M$ ,

- 1 Train predictor  $f_m(\cdot)$  to minimize some weighted loss  $\sum_{i=1}^n w_{mi} L(y_i, f_m(\mathbf{x}_i))$
- 2 Assign larger weights  $w_{m+1,i}$  to samples  $i$  with higher loss

- The final prediction is done using a weighted average

$$F_M(\mathbf{x}) = \sum_{m=1}^M \alpha_m f_m(\mathbf{x}), \quad \text{for some learned weights } \{\alpha_m\}$$

- For simplicity, we will focus on the design of binary classifiers  $f_m(\cdot) \in \{-1, 1\}$ 
  - Other versions of Adaboost exist for regression and non-binary classification

# Adaboost: Derivation for binary classification

- Consider binary classification, with training data  $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$  for  $y_i = \pm 1$
- Adaboost's base classifiers  $f_m(\cdot) = \pm 1$  are trained for  $m = 1 \dots M$  as follows:
  - Define the step- $m$  **boosted classifier**  $F_m(\mathbf{x}) \triangleq F_{m-1}(\mathbf{x}) + \alpha_m f_m(\mathbf{x})$
  - Initialize  $F_0(\mathbf{x}) = 0$ . For  $m = 1, \dots, M$ , design  $\alpha_m$  and  $f_m$  so that  $F_m$  minimizes the **exponential loss**  $\mathcal{L}_m = \sum_{i=1}^n e^{-y_i F_m(\mathbf{x}_i)}$
- Plugging  $F_m(\cdot)$  into  $\mathcal{L}_m$ , we find that

$$\begin{aligned}
 \mathcal{L}_m &= \sum_{i=1}^n e^{-y_i (F_{m-1}(\mathbf{x}_i) + \alpha_m f_m(\mathbf{x}_i))} = \sum_{i=1}^n \overbrace{e^{-y_i F_{m-1}(\mathbf{x}_i)}}^{\triangleq w_{mi}} e^{-y_i \alpha_m f_m(\mathbf{x}_i)} \\
 &= \sum_{i: y_i = f_m(\mathbf{x}_i)} w_{mi} e^{-\alpha_m} + \sum_{i: y_i \neq f_m(\mathbf{x}_i)} w_{mi} e^{\alpha_m} \quad \text{since } f_m \text{ and } y_i \text{ are } \pm 1 \\
 &= \sum_{i=1}^n w_{mi} e^{-\alpha_m} + (e^{\alpha_m} - e^{-\alpha_m}) \sum_{i: y_i \neq f_m(\mathbf{x}_i)} w_{mi}
 \end{aligned}$$

# Adaboost: Derivation for binary classification (cont.)

- From the previous expression, we see that the  $\mathcal{L}_m$ -minimizing classifier  $f_m(\cdot)$  is that which minimizes the **weighted misclassification loss**

$$\sum_{i: y_i \neq f_m(\mathbf{x}_i)} w_{mi} = \sum_{i=1}^n w_{mi} \mathbb{1}_{y_i \neq f_m(\mathbf{x}_i)}$$

- At step  $m=1$ , use  $w_{1i} = 1 \ \forall i$ , and thus train  $f_1(\cdot)$  on  $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$  as usual
- At step  $m > 1$ , the **weight**  $w_{mi}$  is large if  $F_{m-1}$  made a wrong decision on  $y_i$
- The  $\mathcal{L}_m$ -minimizing  $\alpha_m$  can be found by solving  $\partial \mathcal{L}_m / \partial \alpha_m = 0$ , which gives


$$\alpha_m = \frac{1}{2} \left( \frac{1 - \epsilon_m}{\epsilon_m} \right) \text{ for } \text{weighted error rate } \epsilon_m = \frac{\sum_{i: y_i \neq f_m(\mathbf{x}_i)} w_{mi}}{\sum_i w_{mi}}$$

- In summary, to train Adaboost, initialize  $w_{1i} = 1 \ \forall i$  and then, for  $m = 1 \dots M$ , compute  $f_m(\cdot)$ ,  $\epsilon_m$ ,  $\alpha_m$ ,  $F_m(\cdot)$ , and  $\{w_{m+1,i}\}_{i=1}^n$
- Once trained, predict the label of a test sample  $\mathbf{x}$  using  $F_M(\mathbf{x})$

# Adaboost: Implementation

- Adaboost classification implemented via `sklearn.AdaBoostClassifier`
  - By default, it uses 50 decision tree classifiers of depth 1 ("stumps")
- Adaboost regression implemented via `sklearn.AdaBoostRegressor`
  - By default, it uses 50 decision tree regressors of depth 3

*AdaBoost Classifier Working Principle with Decision Stump as a Base Classifier*



# Gradient boosting: Intuitions

- **Gradient boosting** is a generalization / re-interpretation of Adaboost
- For intuition, suppose we want to learn a weighted ensemble  $F_M(\mathbf{x}) = \sum_{m=1}^M \alpha_m f_m(\mathbf{x})$  that minimizes the RSS  $\mathcal{L} = \sum_{i=1}^n (y_i - F_M(\mathbf{x}_i))^2$
- Like with Adaboost, the idea is to train  $\{f_m(\cdot)\}$  sequentially
  - At step  $m = 1$ , we train  $f_1(\cdot)$  to minimize RSS on  $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$  as usual
  - At step  $m = 2, \dots, M$ , we construct  $F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \alpha_m f_m(\mathbf{x})$
  - But how do we train  $\alpha_m$  and  $f_m(\cdot) \in \mathcal{F}$ ?
- Intuition: A “perfect” predictor  $F_m(\cdot)$  would yield zero error, i.e.,
 
$$\forall i : y_i = F_m(\mathbf{x}_i) \quad \Leftrightarrow \quad \forall i : \alpha_m f_m(\mathbf{x}_i) = y_i - F_{m-1}(\mathbf{x}_i) \triangleq r_{mi}$$
  - So  $\alpha_m f_m(\cdot)$  should predict the **residual error**  $r_{mi}$  of the previous step!
- Note that the residual  $y - F$  is the negative gradient of  $\frac{1}{2}(y - F)^2$  w.r.t.  $F$ 
  - Suggests a way to generalize from RSS cost to a generic cost ...

# Gradient boosting: Approach

- Goal: Design predictor  $F_M(\cdot)$  to minimize a training loss of the form

$$\mathcal{L} = \sum_{i=1}^n L(y_i, F_M(\mathbf{x}_i)) \quad \text{for some } L(\cdot, \cdot)$$

where  $F_M(\cdot)$  is sequentially learned from base predictors  $f_m(\cdot) \in \mathcal{F}$

- Adaboost did this with exponential loss  $L(y, F) = e^{-yF}$  and binary  $f_m = \pm 1$
- What about general  $L$  and  $\mathcal{F}$ ?
- Ideally, we would like to do the following:
  - Train  $f_1 \in \mathcal{F}$  to minimize  $\sum_{i=1}^n L(y_i, f_1(\mathbf{x}_i))$ , set  $F_1 = f_1$
  - For  $m = 2 \dots M$ :  $F_m = F_{m-1} + \arg \min_{\alpha_m, f_m \in \mathcal{F}} \sum_{i=1}^n L(y_i, F_{m-1}(\mathbf{x}_i) + \alpha_m f_m(\mathbf{x}_i))$
  - But this optimization is too difficult!
- Idea: Instead of exact optimization, settle for a **gradient step**, i.e.,

$$\text{For } m = 2 \dots M : \quad F_m = F_{m-1} - \alpha_m \nabla \mathcal{L}(F_{m-1})$$



# Gradient boosting: Details

## ■ Problem:

- We are limited to updates of the form

$$F_m = F_{m-1} + \alpha_m f_m, \text{ for } f_m \in \mathcal{F} \text{ where } \mathcal{F} \text{ is set of base predictors}$$

- But gradient descent does not constrain the update to  $\mathcal{F}$ :

$$F_m = F_{m-1} + \alpha_m [-\nabla \mathcal{L}(F_{m-1})], \text{ with } \nabla \mathcal{L}(F_{m-1}) = \sum_{i=1}^n \underbrace{\frac{\partial L(y_i, F_{m-1}(\mathbf{x}_i))}{\partial F}}_{\triangleq -r_{mi}}$$

- Solution: Train  $f_m \in \mathcal{F}$  to be *close to*  $-\nabla \mathcal{L}(F_{m-1})$

- In particular, for  $m = 2, \dots, M$ , we do the following:

- First train  $f_m \in \mathcal{F}$  to minimize  $\sum_{i=1}^n (r_{mi} - f_m(\mathbf{x}_i))^2$

- Then choose  $\alpha_m$  via line-search:

$$\alpha_m = \arg \min_{\alpha} \sum_{i=1}^n L(y_i, F_{m-1}(\mathbf{x}_i) + \alpha f_m(\mathbf{x}_i))$$

- Finally, update boosted predictor:  $F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \alpha_m f_m(\mathbf{x})$

# Gradient boosting: Implementation


- In practice, it helps to slow down the updates using “learning rate”  $\mu \in (0, 1]$ :

$$F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \mu \alpha_m f_m(\mathbf{x})$$

- Classification implemented via `sklearn.GradientBoostingClassifier`
  - By default, it uses 100 decision tree classifiers of depth 3
  - By default, uses  $\mu = 0.1$
- Regression implemented via `sklearn.GradientBoostingRegressor`
  - By default, it uses 100 decision tree regressors of depth 3
  - By default, uses  $\mu = 0.1$

# Gradient boosting: Example

```
from sklearn.ensemble import GradientBoostingRegressor
est = GradientBoostingRegressor(n_estimators=2000, max_depth=1).fit(X, y)
for pred in est.staged_predict(X):
    plt.plot(X[:, 0], pred, color='r', alpha=0.1)
```



from Gradient Boosted Regression Trees by Prettenhofer & Louppe

# Extreme gradient boosting (XGBoost)

- **XGBoost** is an evolution of gradient boosting that has won many recent machine-learning contests
  - For low-dimensional problems, XGBoost often performs as good as neural nets, but it is much easier to design/train!
- Compared to gradient boosting, XGBoost uses...
  - second-order Taylor series approximation
  - regularization
  - sophisticated tree-search
  - many parallelization and hardware acceleration tricks
- The **XGBoost package** is not part of **sklearn**, but compatible with it
  - Once you download the XGBoost package, you can **use it just like sklearn**

# XGBoost: Derivation

- **XGBoost** sequentially designs a boosted predictor  $F_M(\mathbf{x}) = \sum_{m=1}^M f_m(\mathbf{x})$
- Similar to gradient boosting, we'd ideally like to do the following:
  - Train  $f_1 \in \mathcal{F}$  to minimize  $\sum_{i=1}^n L(y_i, f_1(\mathbf{x}_i))$ , set  $F_1 = f_1$
  - For  $m = 2 \dots M$ :  $F_m = F_{m-1} + \arg \min_{f_m \in \mathcal{F}} \sum_{i=1}^n L(y_i, F_{m-1}(\mathbf{x}_i) + f_m(\mathbf{x}_i)) + \phi(f_m)$
  - Note that we added **regularization** to the cost function
  - Unfortunately, this optimization problem is too difficult to solve exactly
- Idea: Simplify using a **2nd-order Taylor series approximation**
  - Use  $L(y_i, F_{m-1}(\mathbf{x}_i) + f_m(\mathbf{x}_i)) \approx L(y_i, F_{m-1}(\mathbf{x}_i)) + g_{mi}f_m(\mathbf{x}_i) + \frac{1}{2}h_{mi}f_m^2(\mathbf{x}_i)$   
 where  $g_{mi} \triangleq \frac{\partial L(y_i, F_{m-1}(\mathbf{x}_i))}{\partial F}$  and  $h_{mi} \triangleq \frac{\partial^2 L(y_i, F_{m-1}(\mathbf{x}_i))}{\partial^2 F}$
  - Then the optimization problem becomes

$$\arg \min_{f_m \in \mathcal{F}} \sum_{i=1}^n [g_{mi}f_m(\mathbf{x}_i) + \frac{1}{2}h_{mi}f_m^2(\mathbf{x}_i)] + \phi(f_m) \quad (1)$$

# XGBoost: Derivation 2

- Now suppose that  $f(\cdot)$  is a **decision tree with  $L$  leaves** (recall page 10)
  - Then we can write  $f(\mathbf{x}) = z_{q(\mathbf{x})}$
  - The function  $q(\mathbf{x}) : \mathbb{R}^d \rightarrow \{1, \dots, L\}$  assigns  $\mathbf{x}$  to a leaf
  - $z_\ell \in \mathbb{R}$  is the predictor output for leaf  $\ell$

- XGBoost uses the regularization

$$\phi(f) = \frac{\lambda}{2} \sum_{\ell=1}^L z_\ell^2 + \gamma L, \quad \text{for some } \lambda > 0 \text{ and } \gamma > 0$$

- With  $f_m(\cdot)$  and  $\phi(f_m)$  as above, the optimization problem (1) becomes

$$\begin{aligned} & \arg \min_{q, \mathbf{z}} \sum_{\ell=1}^L \sum_{i: q(\mathbf{x}_i) = \ell} \left( g_i z_\ell + \frac{1}{2} h_i z_\ell^2 \right) + \frac{\lambda}{2} \sum_{\ell=1}^L z_\ell^2 + \gamma L \\ &= \arg \min_{q, \mathbf{z}} \sum_{\ell=1}^L \left[ \left( \sum_{i: q(\mathbf{x}_i) = \ell} g_i \right) z_\ell + \frac{1}{2} \left( \lambda + \sum_{i: q(\mathbf{x}_i) = \ell} h_i \right) z_\ell^2 + \gamma \right] \end{aligned} \quad (2)$$

## XGBoost: Derivation 3

- We can optimize over  $z \in \mathbb{R}^L$  to yield

$$z_\ell = -\frac{\sum_{i:q(\mathbf{x}_i)=\ell} g_i}{\lambda + \sum_{i:q(\mathbf{x}_i)=\ell} h_i}, \quad \ell = 1, \dots, L$$

and plug this back into (2) to obtain an optimization problem over  $q(\cdot)$ :

$$\arg \min_q \sum_{\ell=1}^L \left[ 2\gamma - \frac{(\sum_{i:q(\mathbf{x}_i)=\ell} g_i)^2}{\lambda + \sum_{i:q(\mathbf{x}_i)=\ell} h_i} \right]$$

- Using the above loss, the tree  $q(\cdot)$  can be designed as on page 11: Choose a feature  $j$  and search for the threshold  $t$  that maximally reduces loss after a split
- For example, consider the split  $S \rightarrow (S_1, S_2)$ . The loss reduction would be

$$-\frac{(\sum_{i \in S} g_i)^2}{\lambda + \sum_{i \in S} h_i} - \left[ 2\gamma - \frac{(\sum_{i \in S_1} g_i)^2}{\lambda + \sum_{i \in S_1} h_i} - \frac{(\sum_{i \in S_2} g_i)^2}{\lambda + \sum_{i \in S_2} h_i} \right]$$

If the maximum loss reduction is negative, then it's better *not* to split!

- Further details can be found in [the original paper](#)

# XGBoost: Example

```

1 # First XGBoost model for Pima Indians dataset
2 from numpy import loadtxt
3 from xgboost import XGBClassifier
4 from sklearn.model_selection import train_test_split
5 from sklearn.metrics import accuracy_score
6 # load data
7 dataset = loadtxt('pima-indians-diabetes.csv', delimiter=",")
8 # split data into X and y
9 X = dataset[:,0:8]
10 Y = dataset[:,8]
11 # split data into train and test sets
12 seed = 7
13 test_size = 0.33
14 X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=test_size, ran
15 # fit model no training data
16 model = XGBClassifier()
17 model.fit(X_train, y_train)
18 # make predictions for test data
19 y_pred = model.predict(X_test)
20 predictions = [round(value) for value in y_pred]
21 # evaluate predictions
22 accuracy = accuracy_score(y_test, predictions)
23 print("Accuracy: %.2f%%" % (accuracy * 100.0))

```

Running this example produces the following output.

```

1 Accuracy: 77.95%

```

This is a [good accuracy score on this problem](#), which we would expect, given the capabilities of the model and the modest complexity of the problem.

[from develop-first-xgboost-model-python-scikit-learn](#)



# Learning objectives

- Understand intuition behind ensemble methods: “the wisdom of the crowd”
- Understand parallel ensemble methods
  - bagging, pasting
  - random feature selection
- Understand decision trees
  - feature thresholding and decision regions
  - training via top-down tree induction
  - homogeneity metrics: variance reduction, gini impurity
  - ensemble extension: random forests
- Understand boosting, or sequentially trained ensemble methods
  - Adaboost
  - gradient boosting
  - XGBoost